

Week 1 Internship Report: OUBT Bootcamp – Day 1 to Day 5 Learnings:

NAME: LIKHITH SASANK UPPALAPATI VENKATA

MAIL : likhith2kuv@gmail.com

Introduction

This document summarizes my learnings and activities during Week 1 of the OUBT Bootcamp Internship (Day 1 to Day 5). Over the course of the week, I focused on building foundational skills in AWS services, Python for data engineering, SQL, and GitHub version control. I also gained hands-on experience with **Git and GitHub**, including repository management, branching, and committing changes, which helped me track and manage my code effectively. Each day involved hands-on exercises, practical tasks, and mini-projects aimed at developing core competencies in cloud computing, data processing, and automation workflows. The following sections detail the tasks completed, concepts learned, and technical skills acquired throughout the week.

Project Repository Details:

All Week 1 code, scripts, and documentation have been maintained in my Git repository: https://github.com/likhith5697/Week1_OUBT_Learnings. The repository includes Python scripts, SQL queries, and AWS configuration files for tasks completed from Day 1 to Day 5.

**NAME: LIKHITH SASANK UPPALAPATI
VENKATA**

DAY: 1

AWS IAM Setup:

Create IAM users, groups, and roles.

Assign permissions and policies.

Test login using IAM sign-in link.

Billing Alert and CloudWatch Alarm:

Enable billing alerts in the AWS Billing Console.

Create a CloudWatch alarm to track monthly charges.

Set up an SNS topic to send cost alerts via email.

Amazon S3 Bucket:

Create a new S3 bucket.

Upload files of different types (PDF, image, CSV).

Enable bucket versioning and upload the same file again to check multiple versions.

GitHub Setup:

Create a new repository for OUBT bootcamp tasks.

Clone the repository locally.

Add the Day 1 Word document and screenshots.

Commit and push the files to GitHub.

1. Introduction

This week focused on AWS fundamentals, IAM, S3, GitHub basics, and Python for Data Engineering. The goal was to understand access management, cloud storage, and data handling concepts.

2. IAM Concepts Summary

IAM Overview:

IAM – It is a security's service that will control who can access our AWS account

IAM – USERS:

Users means one individual person that needs an access

Example : Let's say we are a team of developers, admin, testers. Each person needs a particular access so we create users as Likhith.Dev, Sasank.Admin, Venkat.Tester and give them roles and permissions as shown below.

Likhith.Dev -> **Developer** -> S3, Glue, RDS, Lambda

Sasank.Admin -> **Admin** -> Full Access

Venkat.Tester -> **Tester** -> Read Only

Here we attach policies to Users.

IAM – GROUPS:

Groups are the collection of Users who share same permissions.

Here we do assign policies to groups not for users separately as shown below,

Developers -> Access to S3, Glue, RDS, Lambda -> Members (Likhith, Sasank)

Admin -> Full Access -> Members (Venkat)

Tester -> Read Only Access (Durga)

IAM ROLES:

Role is a temporary identity with specific permissions.

AWS service assumes a role

Users assume a role

Group assumes a role

Example : Suppose I have a microservice or may be a spring boot app deployed in EC2 that basically needs to read and write to RDS, So instead of hardcoding credentials, I can create a role and assign to our EC2 instance for access to RDS, So it gets permission automatically.

IAM Policy:

Policy defines permissions.

It is a json object that says what is allowed and what is not allowed.

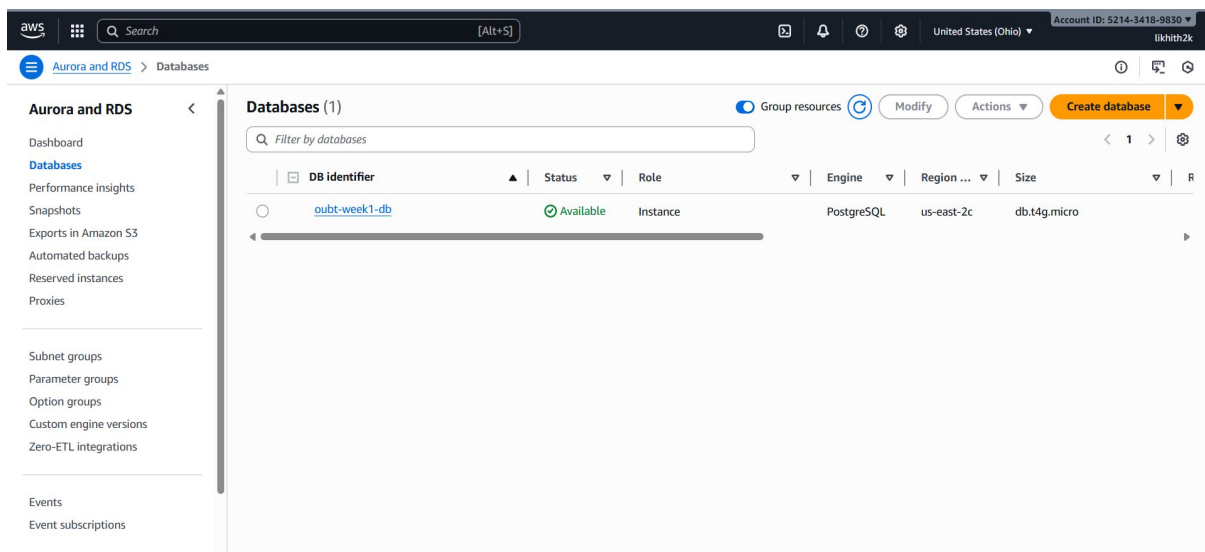
Billing Alerts:

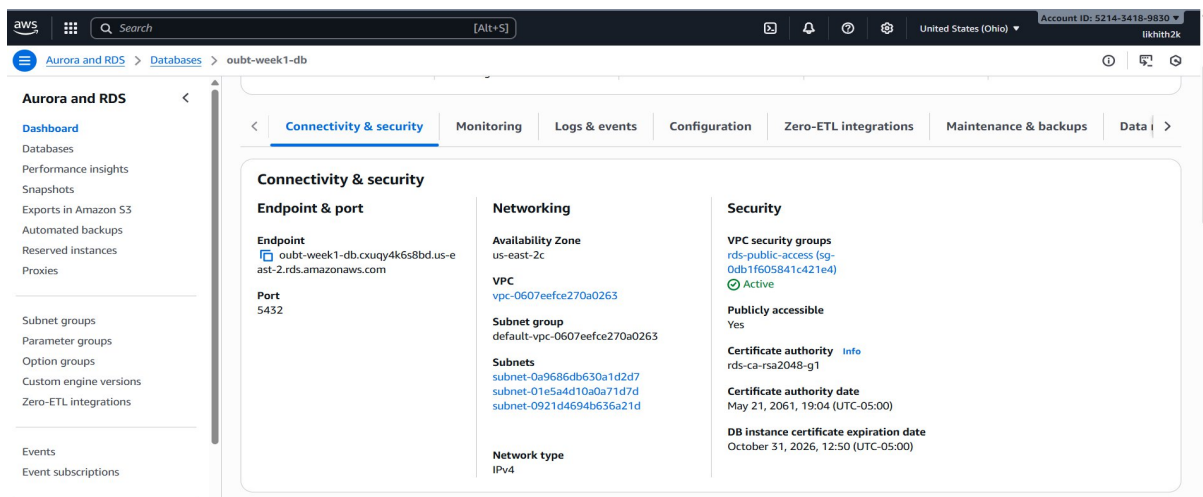
This helps to avoid surprise charges .

We can use cloudwatch alerts / Billing alarms to notify us when usage exceeds a limit.

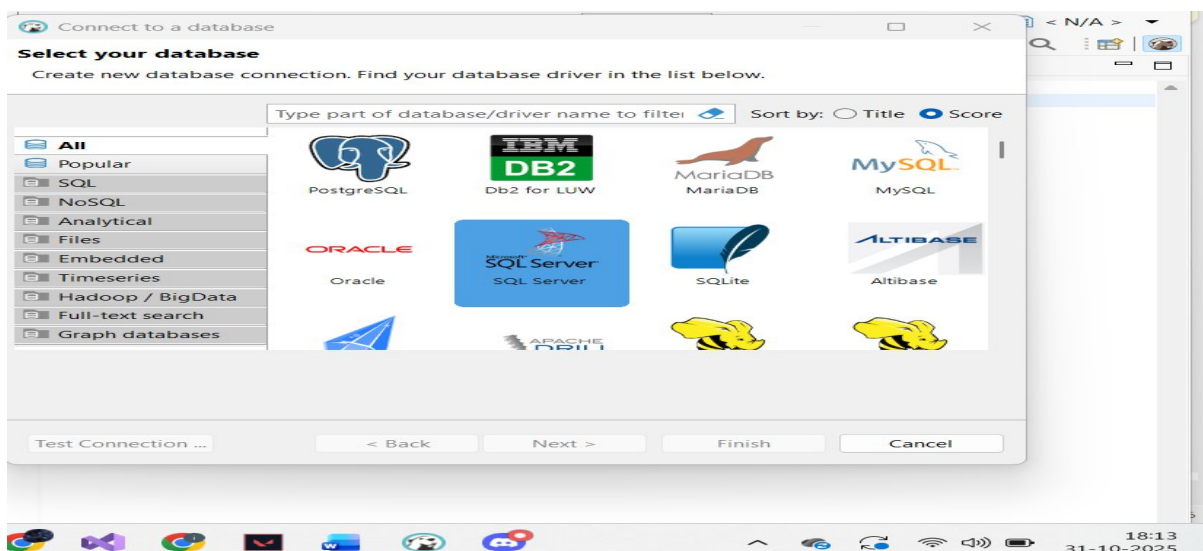
Here Am showing all my hand-ON PRACTICE,

Task 1: Created IAM User

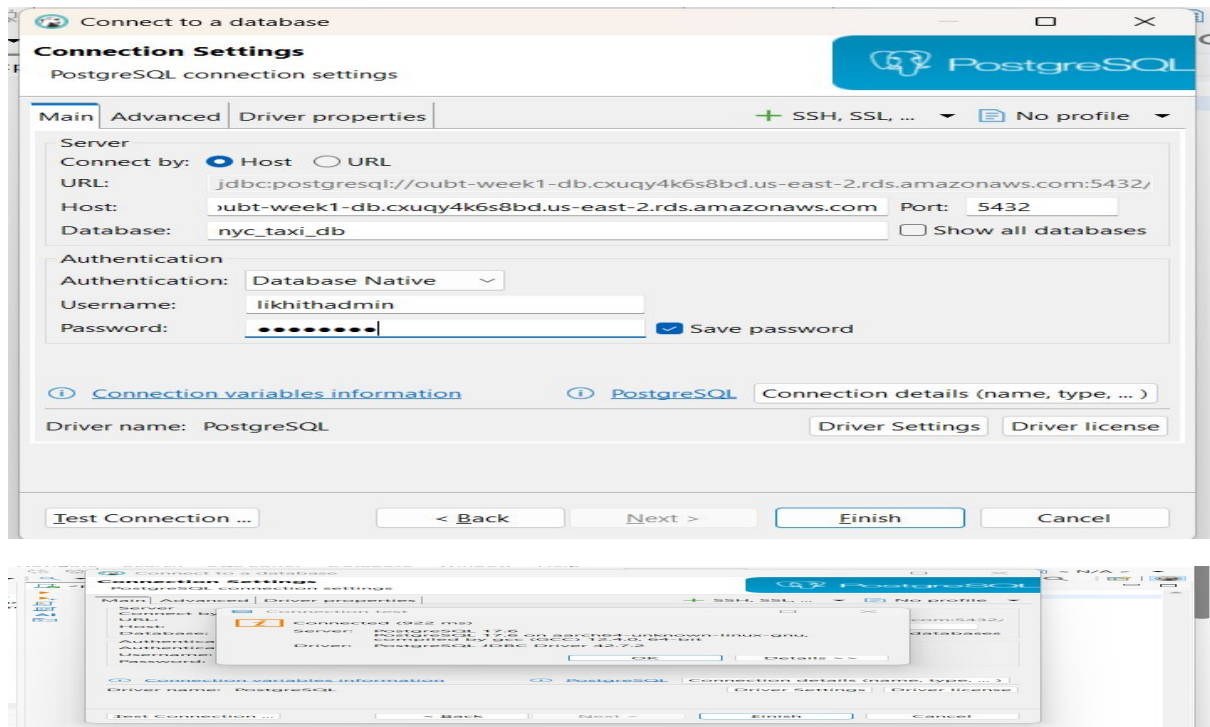




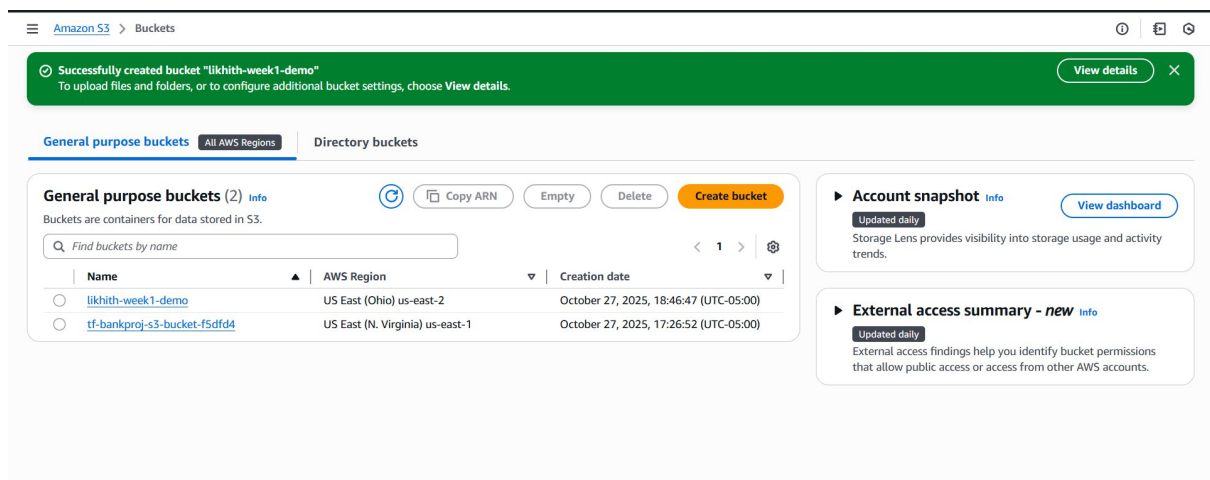
Successfully created user,



Task 2: Created S3 Bucket:

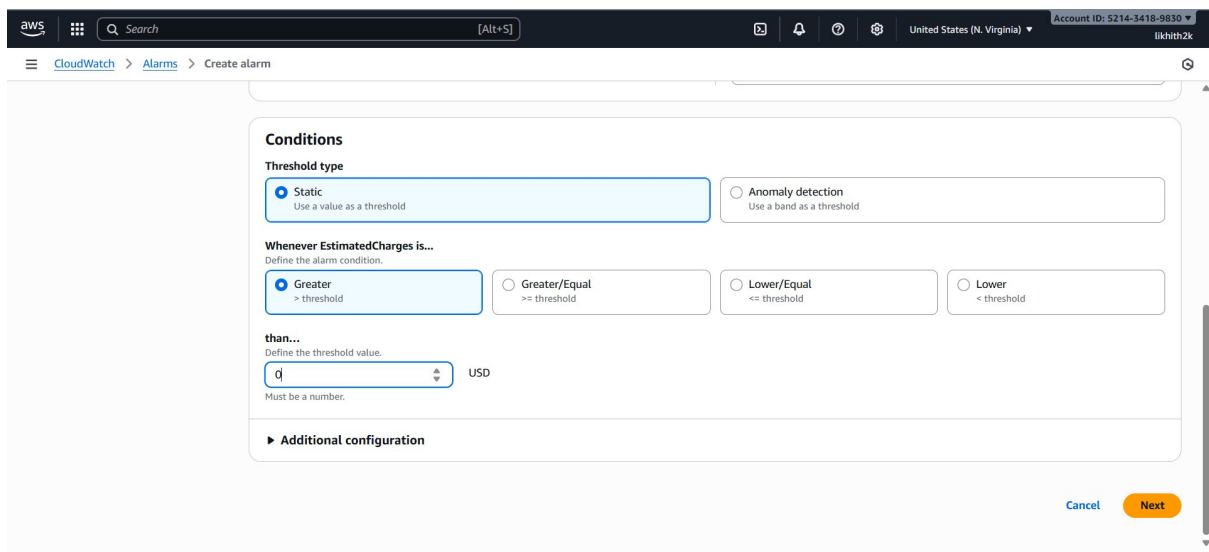
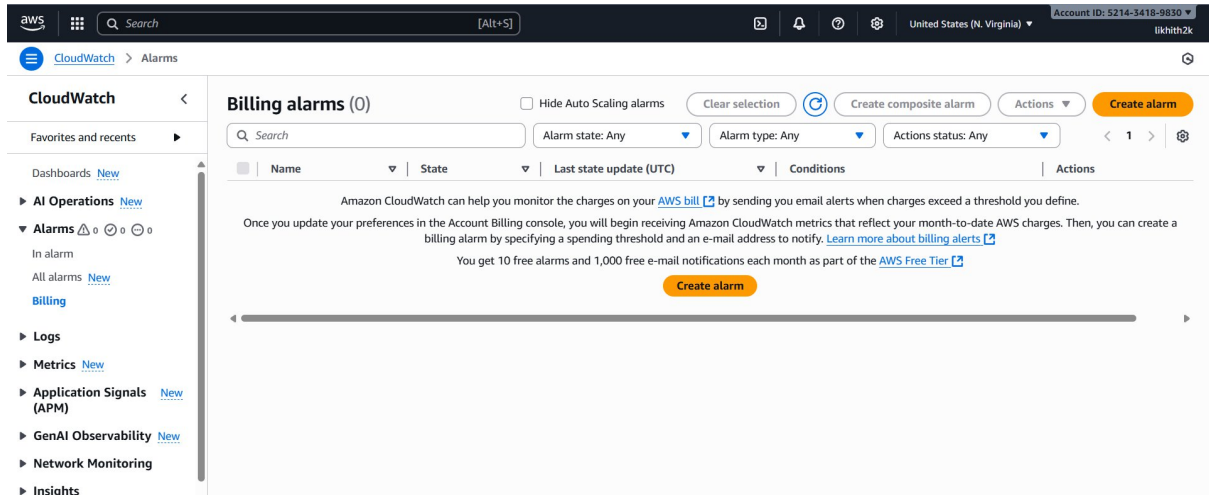


Successfully created S3-BUCKET,



Task 3: Billing Alert

Set up CloudWatch Billing Alarm for \$1 threshold with email alert.



Configure actions

Step 3

Add alarm details

Step 4

Preview and create

Notification

Alarm state trigger
Define the alarm state that will trigger this action.

☒ **In alarm**
The metric or expression is outside of the defined threshold.

☐ **OK**
The metric or expression is within the defined threshold.

☐ **Insufficient data**
The alarm has just started or not enough data is available.

Remove

Send a notification to the following SNS topic
Define the SNS (Simple Notification Service) topic that will receive the notification.

☐ Select an existing SNS topic

☒ Create new topic

☐ Use topic ARN to notify other accounts

Create a new topic...
The topic name must be unique.

SNS topic names can contain only alphanumeric characters, hyphens (-) and underscores (_).

Email endpoints that will receive the notification...
Add a comma-separated list of email addresses. Each address will be added as a subscription to the topic above.

user1@example.com, user2@example.com

Create topic

Add notification

Finally created an alarm,

aws

Search

[Alt+S]

United States (N. Virginia)

Account ID: 5214-3418-9630

likhith2k

CloudWatch

Alarms

CloudWatch

Favorites and recents

Dashboards [New](#)

AI Operations [New](#)

Alarms [New](#) [0](#) [0](#) [0](#) [0](#)

In alarm

All alarms [New](#)

Billing

Logs

Metrics [New](#)

Application Signals (APM) [New](#)

Cost Anomaly Detection [New](#)

Successfully created alarm max-cost-usage-alert.

View alarm

×

Some subscriptions are pending confirmation

Amazon SNS doesn't send messages to an endpoint until the subscription is confirmed

View SNS Subscriptions

×

Billing alarms (1)

☐ Hide Auto Scaling alarms

Clear selection

Create composite alarm

Actions

Create alarm

Alarm state: Any

Alarm type: Any

Actions status: Any

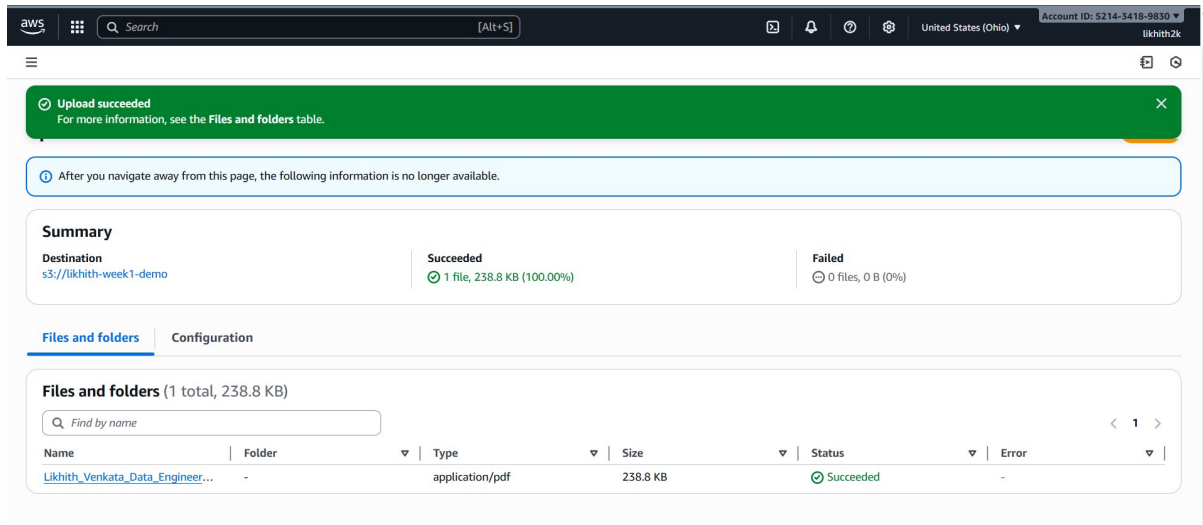
< 1 >

⚙️

<input type="checkbox"/>	Name	State	Last state update (UTC)	Conditions	Actions
<input type="checkbox"/>	max-cost-usage-alert	Insufficient data	2025-10-27 23:55:18	EstimatedCharges > 0 for 1 datapoints within 6 hours	<div>Actions enabled</div> <div>Warning</div>

S3 - Section

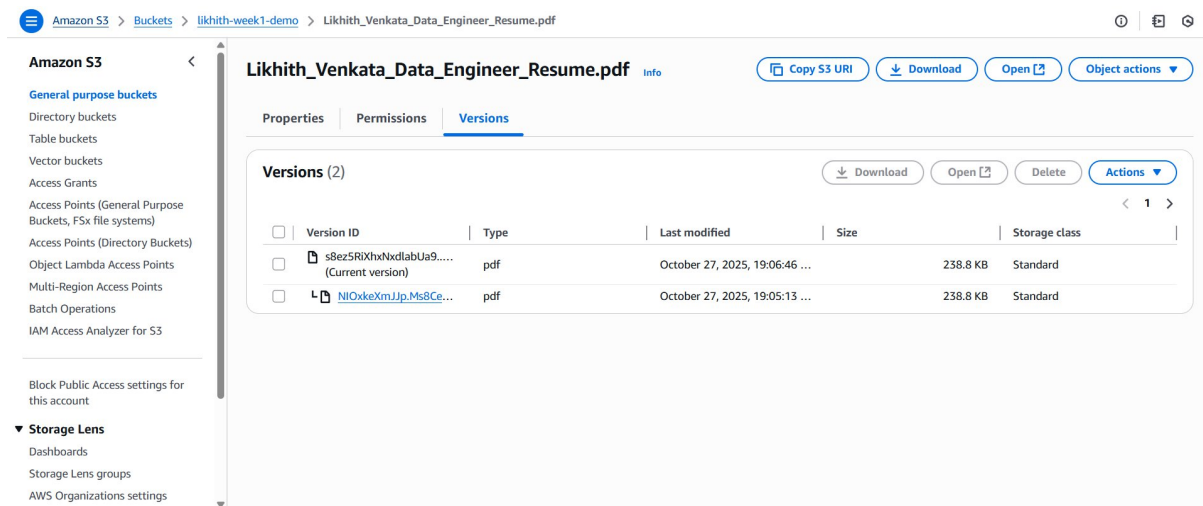
Here I have uploaded the files to S3,



S3 Bucket Versioning Test

I uploaded the same PDF file twice to my S3 bucket. With versioning enabled, both versions of the file are saved. The screenshot below shows the **current version** and the **previous version**. This confirms that versioning is working correctly and older files can be restored if needed.

Screenshot inserted below.



Summary :

Today I learned how AWS **Identity and Access Management (IAM)** helps control and secure user access across AWS services. I created **IAM users, groups, and roles**, assigned appropriate permissions, and understood the concept of least privilege for secure account management.

I also explored **Amazon S3**, where I created a bucket, uploaded multiple file types, and enabled **versioning** to maintain file history and data recovery. In addition, I configured **AWS CloudWatch billing alerts and alarms** to track usage and receive cost notifications, ensuring proactive budget management.

Finally, I set up a **GitHub repository** to store and manage all my bootcamp work, learned the basics of commits and pushing files, and organized my Day 1 deliverables systematically. Overall, I gained hands-on experience in foundational AWS services, account monitoring, and version control — key building blocks for cloud and data engineering.

OUBT Week 1 – Day 2–3: Python & Pandas & Boto3 SDK Practice

Author: Likhith Sasank Uppalapati Venkata

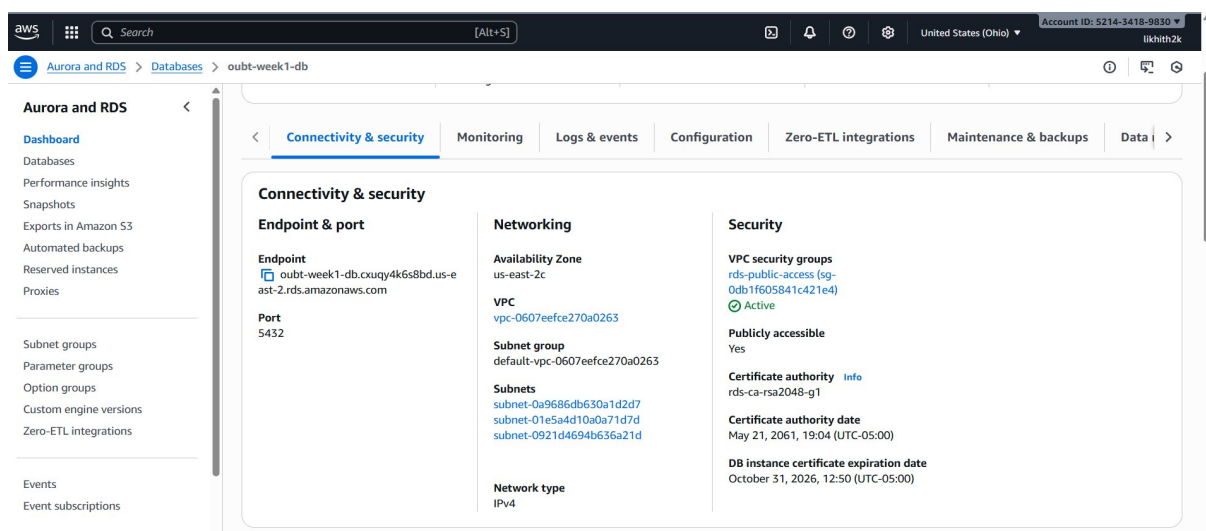
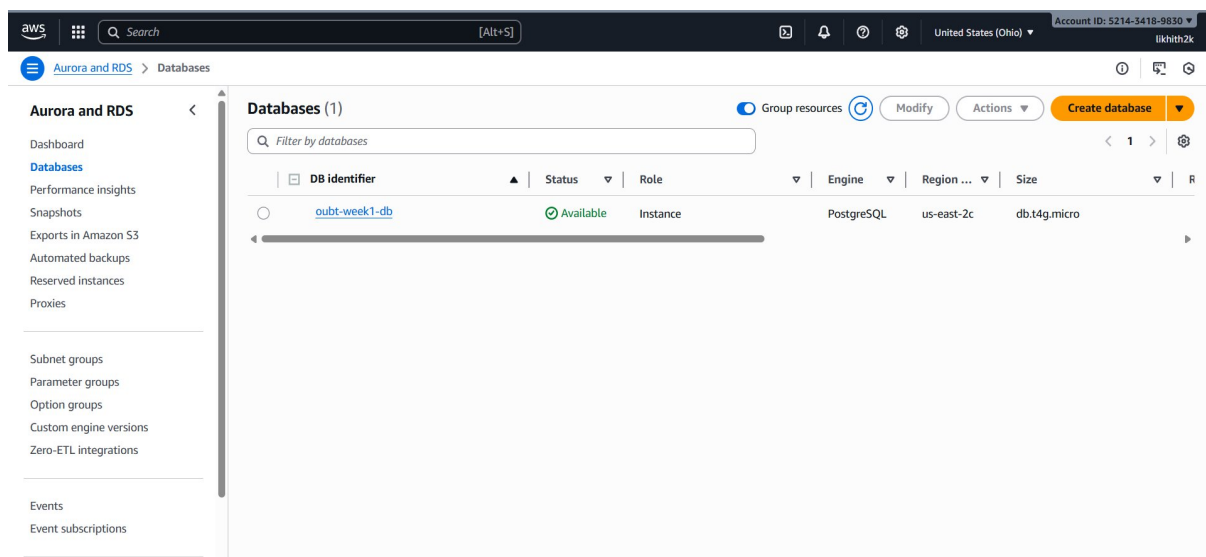
1. Python Practice

Practiced core Python concepts such as lists, dictionaries, loops, functions, and error handling.

Created and managed simple CSV files for student scores and invoice records.

Worked with data filtering and conditional logic using Python structures.

Focused on writing clean, reusable functions and handling exceptions gracefully.





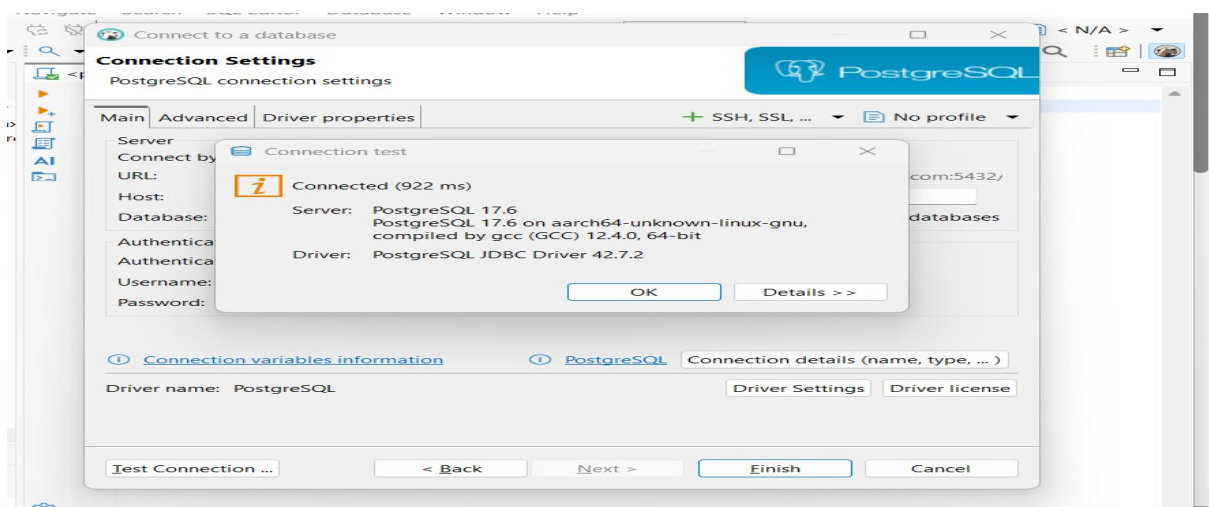
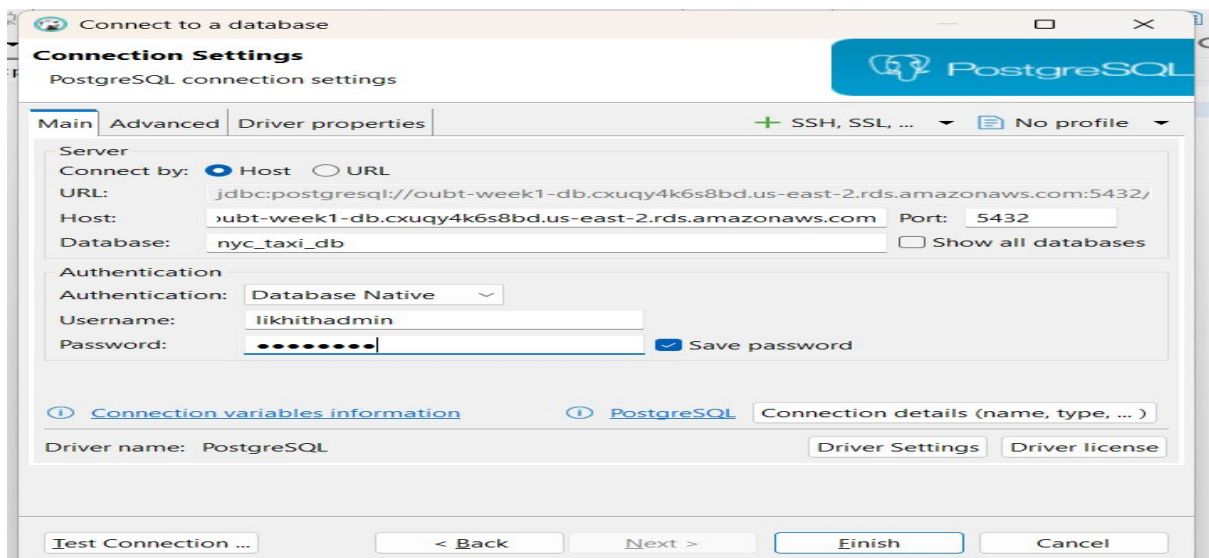
2. Pandas Practice

Learned to create, read, and manipulate datasets using Pandas.

Performed basic data cleaning by removing duplicates and adding calculated columns.

Aggregated and analyzed data to find total sales and quantities.

Combined multiple datasets into a single, cleaned file for further use.



```

63 #Dataframes - merging
64
65 invoice_data = [
66     {"InvoiceId": "I101", "Customer": "Likhith"},
67     {"InvoiceId": "I102", "Customer": "Sasank"}
68 ]
69
70 with open("data/raw/invoice.json", "w") as f:
71     json.dump(invoice_data, f)
72
73
74 invoice_df = pd.read_json("data/raw/invoice.json")
75
76 print("Invoice Data : ")
77 print(invoice_data, "\n")
78
79
80 merged_df = pd.merge(df, invoice_df, on="Customer", how="left")
81 print("Merged Sales and Invoice Data:")
82 print(merged_df.head(), "\n")
83
84 merged_df.to_csv("data/processed/cleaned_sales_data.csv", index=False)
85
86
87 print("Cleaned data saved to data/processed/cleaned_sales_data.csv\n")
88 print("Day 2 Pandas Practice Completed Successfully!")
89
90 sys.stdout.close()

```

3. AWS S3 (Boto3) Practice

Installed and configured Boto3 to connect Python with AWS S3.

Created and managed an S3 bucket for uploading, listing, and downloading files.

```

C:\Users> ikhith > OUBI-TASKS > WEEK1 > Day2_3_Python_Boto3_Pandas_Likhith > scripts > boto3_s3_handson.py > ...
1 import boto3
2 import pandas as pd
3 import sys
4 import os
5 from botocore.exceptions import ClientError
6
7
8 os.makedirs("output", exist_ok=True)
9 sys.stdout = open("output/boto3_s3_output.txt", "w")
10
11 s3 = boto3.client("s3", region_name="us-east-1")
12
13 bucket_name = "likhith-week1-demo"
14 file_path = "data/raw/local_sales.csv"
15 object_name = "sales/local_sales.csv"
16
17 sample_df = pd.DataFrame({
18     "orderId": [1, 2, 3],
19     "Customer": ["Likhith", "Sasank", "Siva"],
20     "Amount": [120, 160, 280]
21 })
22
23 sample_df.to_csv(file_path, index=False)
24
25 print(f"Created local file: {file_path}\n")
26

```

Integrated Pandas with AWS S3 to upload processed data and retrieve it for analysis.

```

28
29 # Here I am sending the files to S3 BUCKET after Creating S3 bucket
30
31 try:
32     s3.create_bucket(Bucket = bucket_name)
33 except ClientError as err:
34     if(err.response["Error"]["Code"] == "BucketAlreadyOwnedByYou"):
35         print(f"Bucket already exists: {bucket_name}")
36     else:
37         print("Error creating bucket:", err)
38
39
40 try:
41     s3.upload_file(file_path, bucket_name, object_name)
42     print(f"Uploaded {file_path} to s3://{bucket_name}/{object_name}")
43 except ClientError as err:
44     print("Upload failed: ", err)
45
46 print("\n Files in S3 bucket:")
47 objects = s3.list_objects_v2(Bucket = bucket_name)
48 if "Contents" in objects:
49     for obj in objects["Contents"]:
50         print(" .", obj["Key"])
51 else:
52     print(" (Bucket is empty) ")
53
54

```

```

# Here I am downloading from s3
download_path = "data/s3_upload/sales_downloaded.csv"
os.makedirs("data/s3_upload", exist_ok=True)

try:
    s3.download_file(bucket_name,object_name,download_path)
    print(f"\n Downloaded file is in {download_path}")
except ClientError as err:
    print("Download failed:", err)

df = pd.read_csv(download_path)
print("\n Data Loaded Back from S3:")
print(df, "\n")

```

Practiced automating simple data workflows between local and cloud storage.

The screenshot shows the AWS Management Console interface for the bucket 'likhith-week1-demo'. The 'Objects' tab is selected, showing a list of objects. The first object is 'Likhith_Venkata_Data_Engine er_Resume.pdf' (238.8 KB, Standard storage class, last modified October 27, 2025). The second object is a folder named 'sales/'. The left sidebar shows the navigation menu with 'Storage Lens' expanded.

Name	Type	Last modified	Size	Storage class
Likhith_Venkata_Data_Engine er_Resume.pdf	pdf	October 27, 2025, 19:06:46 (UTC-05:00)	238.8 KB	Standard
sales/	Folder	-	-	-

All Python Scripts outputs have been saved and organized in a separate folder for reference and documentation purposes.

```
C:\Users> likhi> OUBT-TASKS > WEEK1 > Day2_3_Python_Boto3_Pandas_Likhith > output > boto3_s3_output.txt
1 Created local file: data/raw/local_sales.csv
2
3 Error creating bucket: An error occurred (IllegallocationConstraintException) when calling the CreateBucket operation: The unspecified location constraint is incompatible
4 Uploaded data/raw/local_sales.csv to s3://likhith-week1-demo/sales/local_sales.csv
5
6 Files in S3 bucket:
7 . likhith_Venkata_Data_Engineer_Resume.pdf
8 . sales/local_sales.csv
9
10 Downloaded file is in data/s3_upload/sales_downloaded.csv
11
12 Data Loaded Back from S3:
13   orderId Customer Amount
14 0         1  Likhith  120
15 1         2  Sasank   160
16 2         3   Siva   280
17
18 After doing Tax calculation:
19   orderId Customer Amount Tax TotalWithTax
20 0         1  Likhith  120  12.0      132.0
21 1         2  Sasank   160  16.0      176.0
22 2         3   Siva   280  28.0      308.0
23
24 Final data saved locally in data/s3_upload/sales_with_tax.csv
25 Uploaded final version to s3://likhith-week1-demo/sales/sales_with_tax.csv
26
```

```
C:\Users> likhi> OUBT-TASKS > WEEK1 > Day2_3_Python_Boto3_Pandas_Likhith > output > pandas_basics_output.txt
1 OUBT Week 1 - Day 2-3: Python for Data Engineering
2 Topic: Pandas for Data Manipulation
3
4 Sales Data:   orderId  Customer Product  Quantity Price
5 0         101  Likhith  Laptop         1      800
6 1         102  Sasank   Mobile         2      500
7 2         103  Prathyush Tablet         1      300
8 3         104  Anushka  Laptop         1      800
9 4         105  Sasank   Tablet         3      300
10 5         106  Likhith  Mobile         2      500
11 6         107  Prathyush Laptop         1      800
12 7         108  Anushka  Tablet         2      300
13 8         109  Likhith  Mobile         1      800
14 9         110  Sasank   Laptop         1      500
15
16 Created sample sales_data.csv
17
18 Data read from csv:   orderId  Customer Product  Quantity Price
19 0         101  Likhith  Laptop         1      800
20 1         102  Sasank   Mobile         2      500
21 2         103  Prathyush Tablet         1      300
22 3         104  Anushka  Laptop         1      800
23 4         105  Sasank   Tablet         3      300
24 5         106  Likhith  Mobile         2      500
25 6         107  Prathyush Laptop         1      800
26 7         108  Anushka  Tablet         2      300
27 8         109  Likhith  Mobile         1      800
28 9         110  Sasank   Laptop         1      500
29 |
```



```

71
72 Total Revenue per customer:
73 Customer
74 Anushka      1400
75 Likhith      2600
76 Prathyush    1100
77 Sasank       2400
78 Name: total, dtype: int64
79
80 Total Quantity per Product:
81 Product
82 Laptop       4
83 Mobile       5
84 Tablet       6
85 Name: Quantity, dtype: int64
86
87 Invoice Data :
88 [{'InvoiceId': 'I101', 'Customer': 'Likhith'}, {'InvoiceId': 'I102', 'Customer': 'Sasank'}]
89
90 Merged Sales and Invoice Data:
91 |  orderId  Customer Product  Quantity  Price  total InvoiceId
92 |-----|-----|-----|-----|-----|-----|-----|
93 | 0        101    Likhith  Laptop        1     800     800     I101
94 | 1        102     Sasank  Mobile        2     500    1000     I102
95 | 2        103  Prathyush  Tablet        1     300     300      NaN
96 | 3        104    Anushka  Laptop        1     800     800      NaN
97 | 4        105     Sasank  Tablet        3     300     900     I102
98
99 Cleaned data saved to data/processed/cleaned_sales_data.csv
100 Day 203 Pandas Practice Completed Successfully!
101

```

Amazon S3 > Buckets > likhith-week1-demo > sales/

Amazon S3

General purpose buckets

- Directory buckets
- Table buckets
- Vector buckets
- Access Grants
- Access Points (General Purpose Buckets, FSx file systems)
- Access Points (Directory Buckets)
- Object Lambda Access Points
- Multi-Region Access Points
- Batch Operations
- IAM Access Analyzer for S3

Block Public Access settings for this account

▼ Storage Lens

- Dashboards
- Storage Lens groups

sales/

Copy S3 URI

Objects (1)

Copy S3 URI Copy URL Download Open Delete Actions Create folder Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Show versions

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	local_sales.csv	csv	October 28, 2025, 15:11:43 (UTC-05:00)	66.0 B	Standard

Amazon S3 > Buckets > likhith-week1-demo > sales/

Amazon S3

General purpose buckets

- Directory buckets
- Table buckets
- Vector buckets
- Access Grants
- Access Points (General Purpose Buckets, FSx file systems)
- Access Points (Directory Buckets)
- Object Lambda Access Points
- Multi-Region Access Points
- Batch Operations
- IAM Access Analyzer for S3

Block Public Access settings for this account

▼ Storage Lens

- Dashboards
- Storage Lens groups
- API Operations settings

sales/

Copy S3 URI

Objects (2)

Copy S3 URI Copy URL Download Open Delete Actions Create folder Upload

Find objects by prefix Show versions

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	local_sales.csv	csv	October 28, 2025, 15:18:53 (UTC-05:00)	66.0 B	Standard
<input type="checkbox"/>	sales_with_tax.csv	csv	October 28, 2025, 15:18:53 (UTC-05:00)	116.0 B	Standard

4. Summary

During these sessions, I improved my understanding of Python and file handling, worked on data cleaning and analysis using Pandas, and practiced using AWS S3 with Boto3 for cloud data management. I also organized my work with a proper folder structure, keeping raw, processed, and uploaded data separate, and saving all outputs and screenshots in their own folders. This helped me follow good project practices and build a smooth workflow that connects Python, Pandas, and AWS for managing and analyzing data.

OUBT Week-1 – Day 4: SQL & Data Modeling basics

NAME: LIKHITH SASANK UPPALAPATI VENKATA

Overview:

Day 4 focused on understanding database fundamentals and data modeling concepts. I learned about SQL basics, relational vs. non-relational databases, ACID properties, normalization, and schema evolution. I also studied how to design a proper data model using entities, attributes, and relationships, and explored star schema and Slowly Changing Dimensions (SCD Types 1 & 2). Along with this, I practiced SQL operations like joins, aggregations, and window functions on sample data to strengthen query skills before working with AWS RDS.

Hands-On Practice:

Created sample tables and inserted data for customers, products, and orders.

Practiced SQL operations such as:

Basic filtering using WHERE.

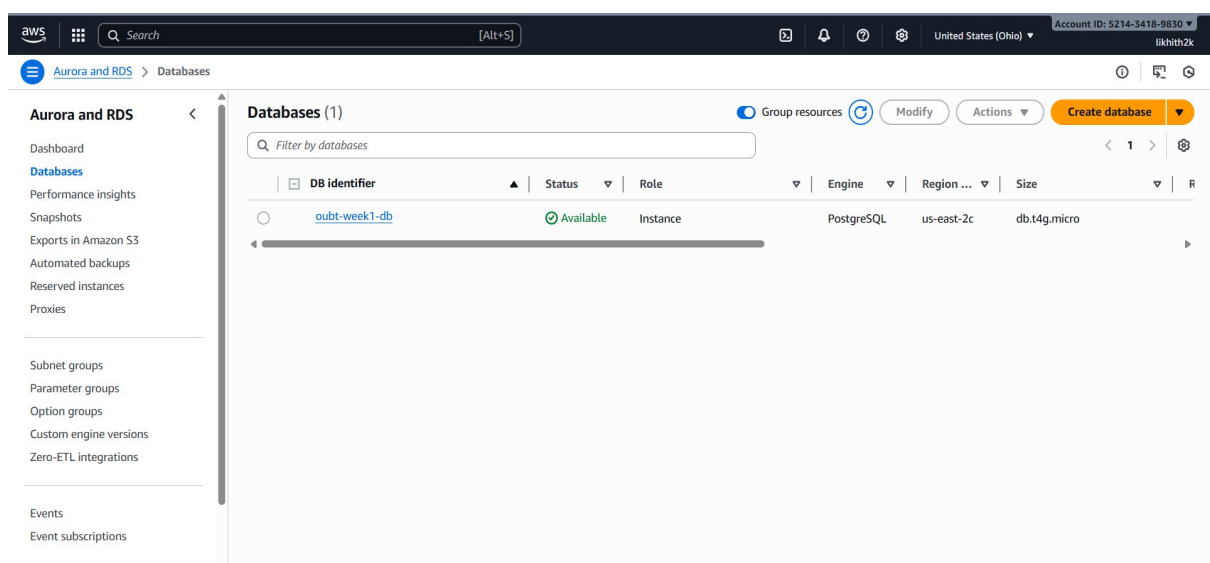
Aggregations with COUNT(), SUM(), and AVG().

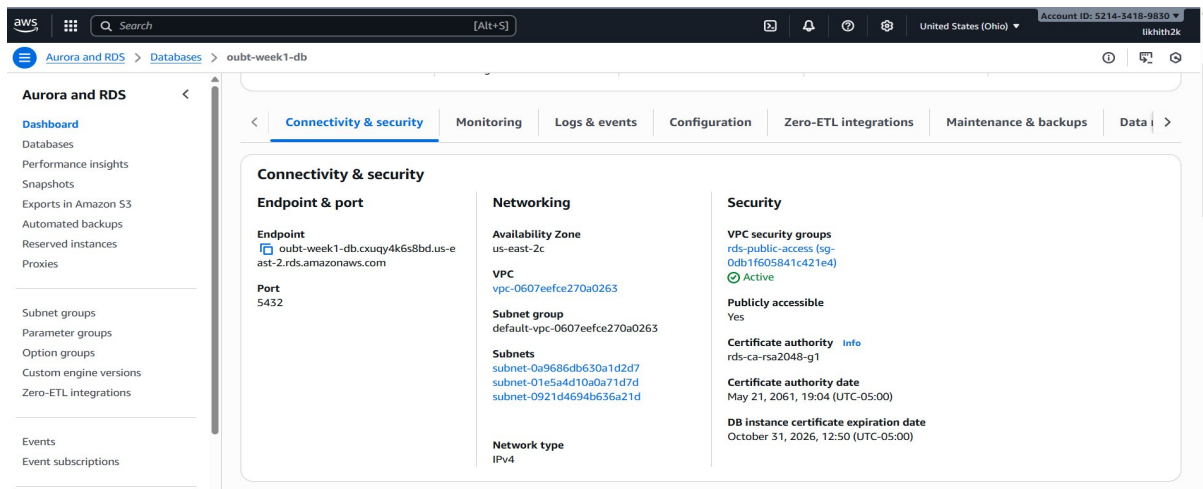
Joins between multiple tables to fetch related information.

Grouping and ordering results.

Used the RANK() window function to rank customers by total spending.

Designed the schema in a **star-schema pattern** with Orders and OrderDetails as fact tables and Customers, Products as dimensions.





SQL – QUERIES:

SELECT * FROM Customers WHERE Country = 'USA';

-- List all products under \$500

SELECT ProductName, Price FROM Products WHERE Price < 500;

-- Count total number of orders

SELECT COUNT(*) AS TotalOrders FROM Orders;

-- Get all orders for customer 'Likhith'

SELECT o.OrderID, o.OrderDate, o.TotalAmount FROM Orders o JOIN Customers c ON
o.CustomerID = c.CustomerID

WHERE c.Name = 'Likhith';

-- Total spent by each customer

```
SELECT c.Name, SUM(o.TotalAmount) AS TotalSpent
FROM Customers c
JOIN Orders o ON c.CustomerID = o.CustomerID
GROUP BY c.Name;
```

-- List all orders with product names

```
SELECT o.OrderID, c.Name AS CustomerName, p.ProductName, od.Quantity
FROM Orders o
JOIN Customers c ON o.CustomerID = c.CustomerID
JOIN OrderDetails od ON o.OrderID = od.OrderID
JOIN Products p ON od.ProductID = p.ProductID;
```

-- Find top 2 customers by spending

```
SELECT c.Name, SUM(o.TotalAmount) AS TotalSpent
FROM Customers c
JOIN Orders o ON c.CustomerID = o.CustomerID
GROUP BY c.Name
ORDER BY TotalSpent DESC
LIMIT 2;
```

-- Find products never ordered

```
SELECT p.ProductName
FROM Products p
LEFT JOIN OrderDetails od ON p.ProductID = od.ProductID
WHERE od.ProductID IS NULL;
```

-- Average order amount per country

```
SELECT c.Country, AVG(o.TotalAmount) AS AvgOrderAmount
```

```
FROM Customers c
```

```
JOIN Orders o ON c.CustomerID = o.CustomerID
```

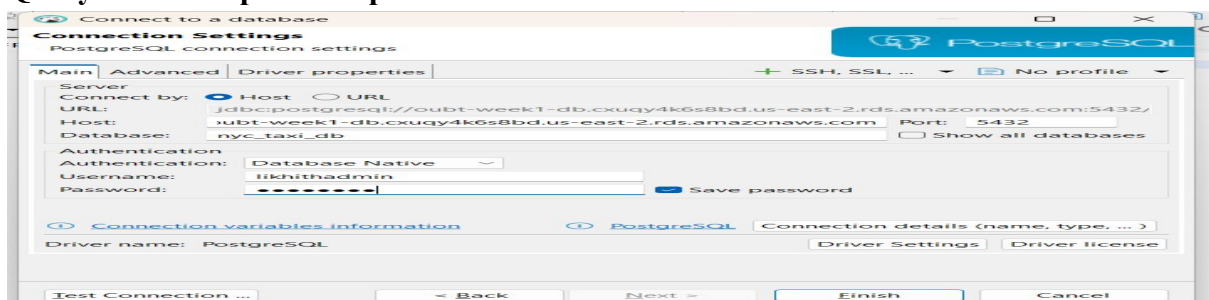
```
GROUP BY c.Country;
```

SQL Query results Screenshots:

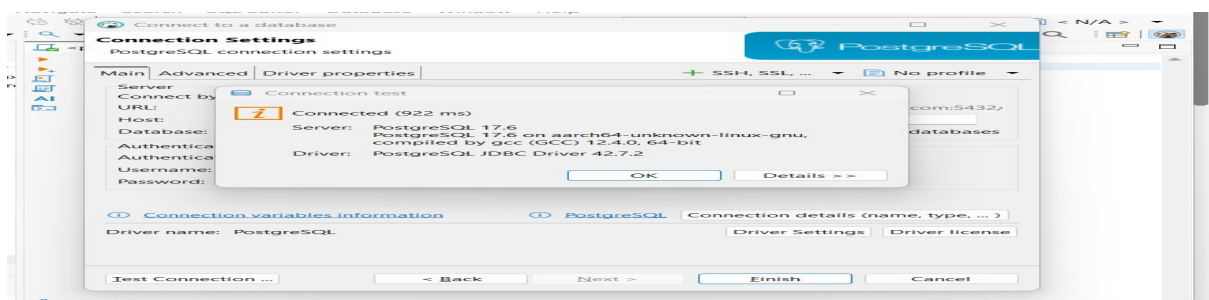
Query 1: Fetch all customers from the USA



Query 2: List all products priced under \$500



Query 3: Count total number of orders



Query 4: Retrieve all orders made by customer “Likhith”

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
OrderID	OrderDate	TotalAmount	
1	2025-10-01	1048.99	

Query 5: Calculate total amount spent by each customer

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
Name	TotalSpent		
Likhith	1048.99		
Sasank	2199.00		
Venkat	1800.00		
Durga	120.00		
Yuvraj	999.00		

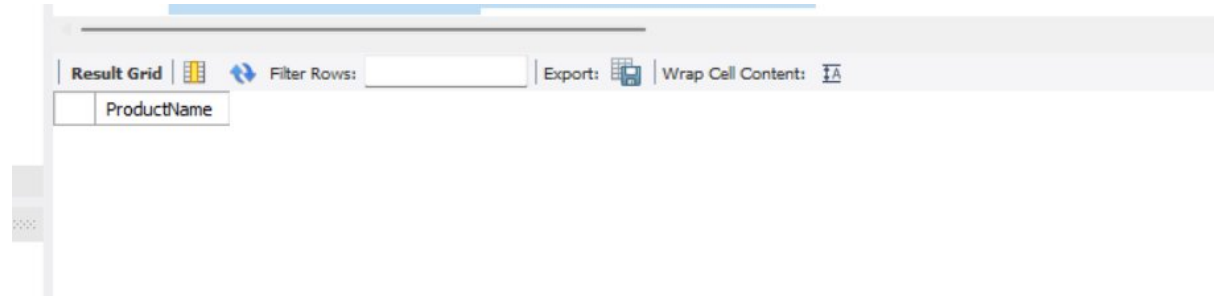
Query 6: Display all orders with product names and quantities

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
OrderID	CustomerName	ProductName	Quantity
1	Likhith	Whey Protein	1
1	Likhith	Jacket	1
2	Sasank	iPad	1
2	Sasank	Ring	2
3	Venkat	Gaming PC	1
4	Durga	Jacket	1
5	Yuvraj	iPad	1

Query 7: Find top 2 customers based on total spending

Result Grid	Filter Rows:	Export:	Wrap Cell Content:	Fetch rows:
Name	TotalSpent			
Sasank	2199.00			
Venkat	1800.00			

Query 8: Identify products that were never ordered

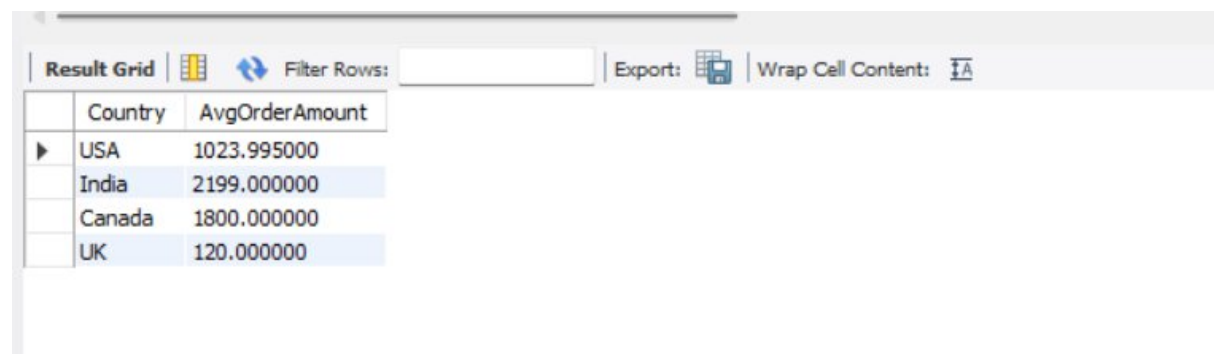


The screenshot shows a database interface with a toolbar at the top containing 'Result Grid', 'Filter Rows', 'Export', and 'Wrap Cell Content'. Below the toolbar is a table with one column labeled 'ProductName'.

ProductName

Query 9: Find average order amount by country

```
4  -- Average order amount per country
5  • SELECT c.Country, AVG(o.TotalAmount) AS AvgOrderAmount
6  FROM Customers c
7  JOIN Orders o ON c.CustomerID = o.CustomerID
8  GROUP BY c.Country;
9
10
```



The screenshot shows a database interface with a toolbar at the top. Below the toolbar is a table with two columns: 'Country' and 'AvgOrderAmount'. The table contains four rows of data.

Country	AvgOrderAmount
USA	1023.995000
India	2199.000000
Canada	1800.000000
UK	120.000000

Key Learnings:

Learned how relational databases store and manage structured data.

Understood **ACID properties**, normalization, and schema design concepts.

Practiced real SQL operations – joins, aggregations, and ranking queries.

Learned how **star schema modeling** supports analytical queries.

Improved understanding of how to represent one-to-many and many-to-many relationships.

Summary:

Overall, Day 4 helped me connect theoretical database concepts with real SQL practice. Building the e-commerce schema gave me a clear idea of how relational data is structured, queried, and optimized. This day built a strong base for moving into AWS RDS and advanced data modeling in the next steps.

OUBT Week-1 – Day 5: SQL & Data Modeling (AWS RDS)

NAME : LIKHITH SASANK UPPALAPATI VENKATA

Overview:

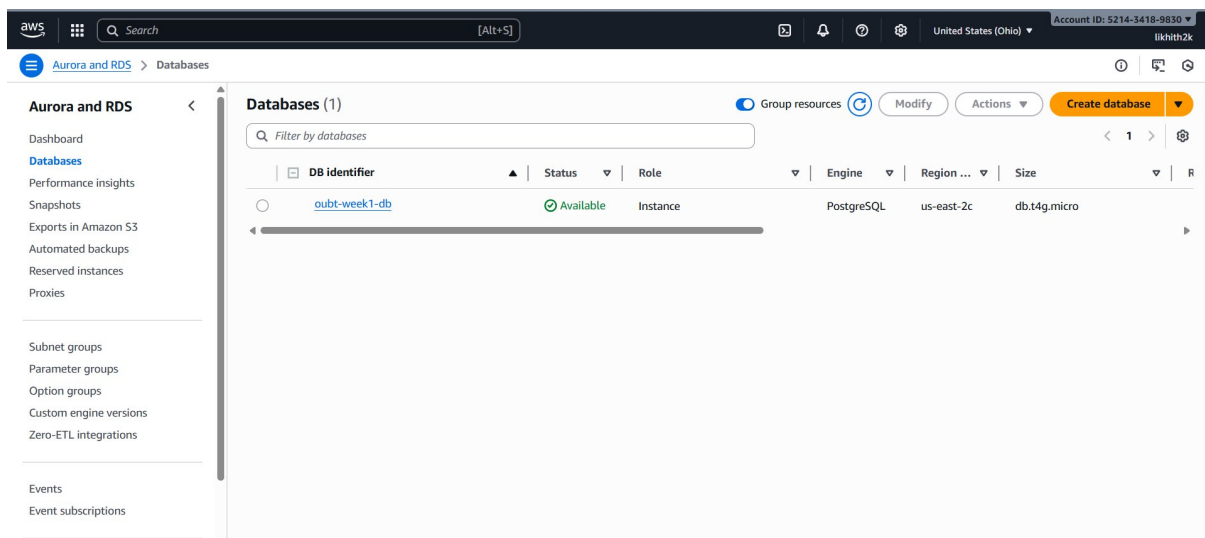
Day 5 focused on setting up and working with a live PostgreSQL database on AWS RDS. I created the instance, configured IAM access, and connected through DBeaver to build tables and run SQL queries on sample taxi data. I also checked RDS metrics in CloudWatch to confirm query activity and understand how CPU and connections change when the database is in use. This helped me connect SQL practice with real AWS database monitoring and management

Objective:

Design a simple relational schema on AWS RDS PostgreSQL and practice SQL operations on sample NYC Taxi Trips data.

Steps Performed:

Created an **AWS RDS PostgreSQL** instance under Free Tier with public access enabled.



Configured inbound rules (port 5432) to allow access from local machine.

The screenshot displays the AWS Management Console interface for an Amazon RDS instance. The top navigation bar includes the AWS logo, a search bar, and account information (United States (Ohio), Account ID: 5214-3418-9830). The left sidebar shows the 'Aurora and RDS' menu with options like Dashboard, Databases, Performance insights, Snapshots, Exports in Amazon S3, Automated backups, Reserved instances, Proxies, Subnet groups, Parameter groups, Option groups, Custom engine versions, Zero-ETL integrations, Events, and Event subscriptions. The main content area is titled 'oubt-week1-db' and features a tabbed interface with 'Connectivity & security' selected. This tab is divided into three sections: 'Endpoint & port', 'Networking', and 'Security'. The 'Endpoint & port' section shows the endpoint 'oubt-week1-db.cxuy4k6s8bd.us-east-2.rds.amazonaws.com' and port '5432'. The 'Networking' section lists the availability zone 'us-east-2c', VPC 'vpc-0607eefce270a0263', subnet group 'default-vpc-0607eefce270a0263', and subnets 'subnet-0a9686db630a1d2d7', 'subnet-01e5a4d10a0a71d7d', and 'subnet-0921d4694b636a21d'. The 'Security' section shows VPC security groups 'rds-public-access (sg-0db1f605841c421e4)' as active, public accessibility set to 'Yes', certificate authority 'rds-ca-rsa2048-g1', certificate authority date 'May 21, 2061, 19:04 (UTC-05:00)', and DB instance certificate expiration date 'October 31, 2026, 12:50 (UTC-05:00)'.

Connectivity & security

Endpoint & port

Endpoint
oubt-week1-db.cxuy4k6s8bd.us-east-2.rds.amazonaws.com

Port
5432

Networking

Availability Zone
us-east-2c

VPC
vpc-0607eefce270a0263

Subnet group
default-vpc-0607eefce270a0263

Subnets
subnet-0a9686db630a1d2d7
subnet-01e5a4d10a0a71d7d
subnet-0921d4694b636a21d

Network type
IPv4

Security

VPC security groups
rds-public-access (sg-0db1f605841c421e4)
Active

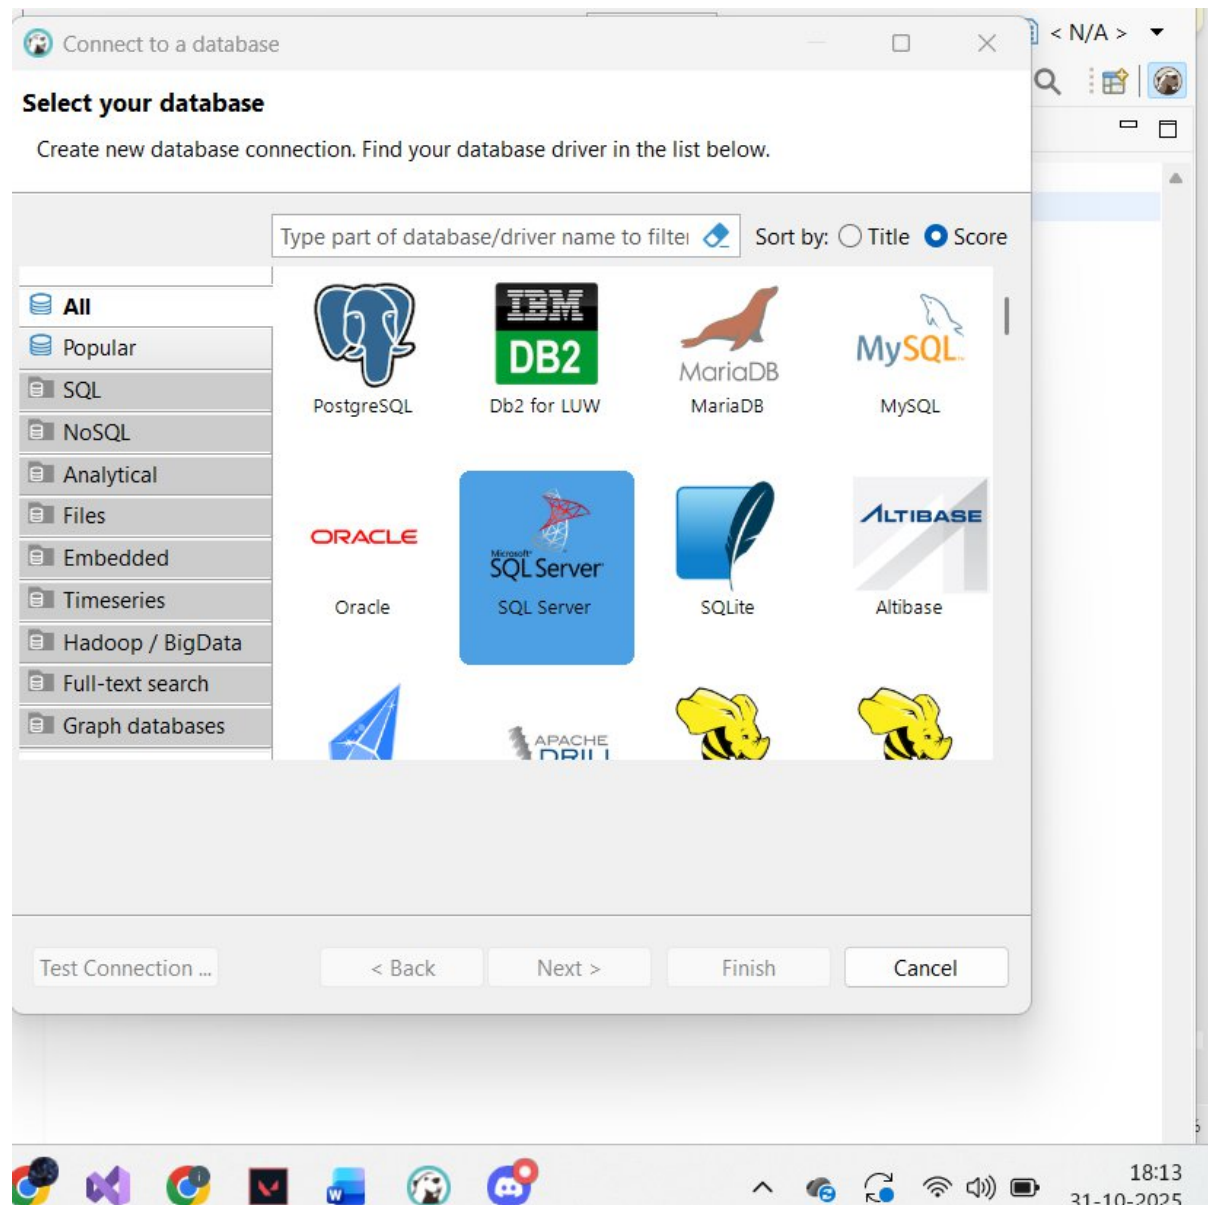
Publicly accessible
Yes

Certificate authority info
rds-ca-rsa2048-g1

Certificate authority date
May 21, 2061, 19:04 (UTC-05:00)

DB instance certificate expiration date
October 31, 2026, 12:50 (UTC-05:00)


Connected to the RDS instance through **DBeaver** using the database endpoint.



Connect to a database

Connection Settings

PostgreSQL connection settings

 PostgreSQL

Main

Advanced

Driver properties

+ SSH, SSL, ...

No profile

Server

Connect by: ☒ Host ☐ URL

URL: jdbc:postgresql://oubt-week1-db.cxuqy4k6s8bd.us-east-2.rds.amazonaws.com:5432/

Host: oubt-week1-db.cxuqy4k6s8bd.us-east-2.rds.amazonaws.com Port: 5432

Database: nyc_taxi_db ☐ Show all databases

Authentication

Authentication: Database Native

Username: likhithadmin

Password:

.....

☒ Save password

Connection variables information

PostgreSQL

Connection details (name, type, ...)

Driver name: PostgreSQL

Driver Settings

Driver license

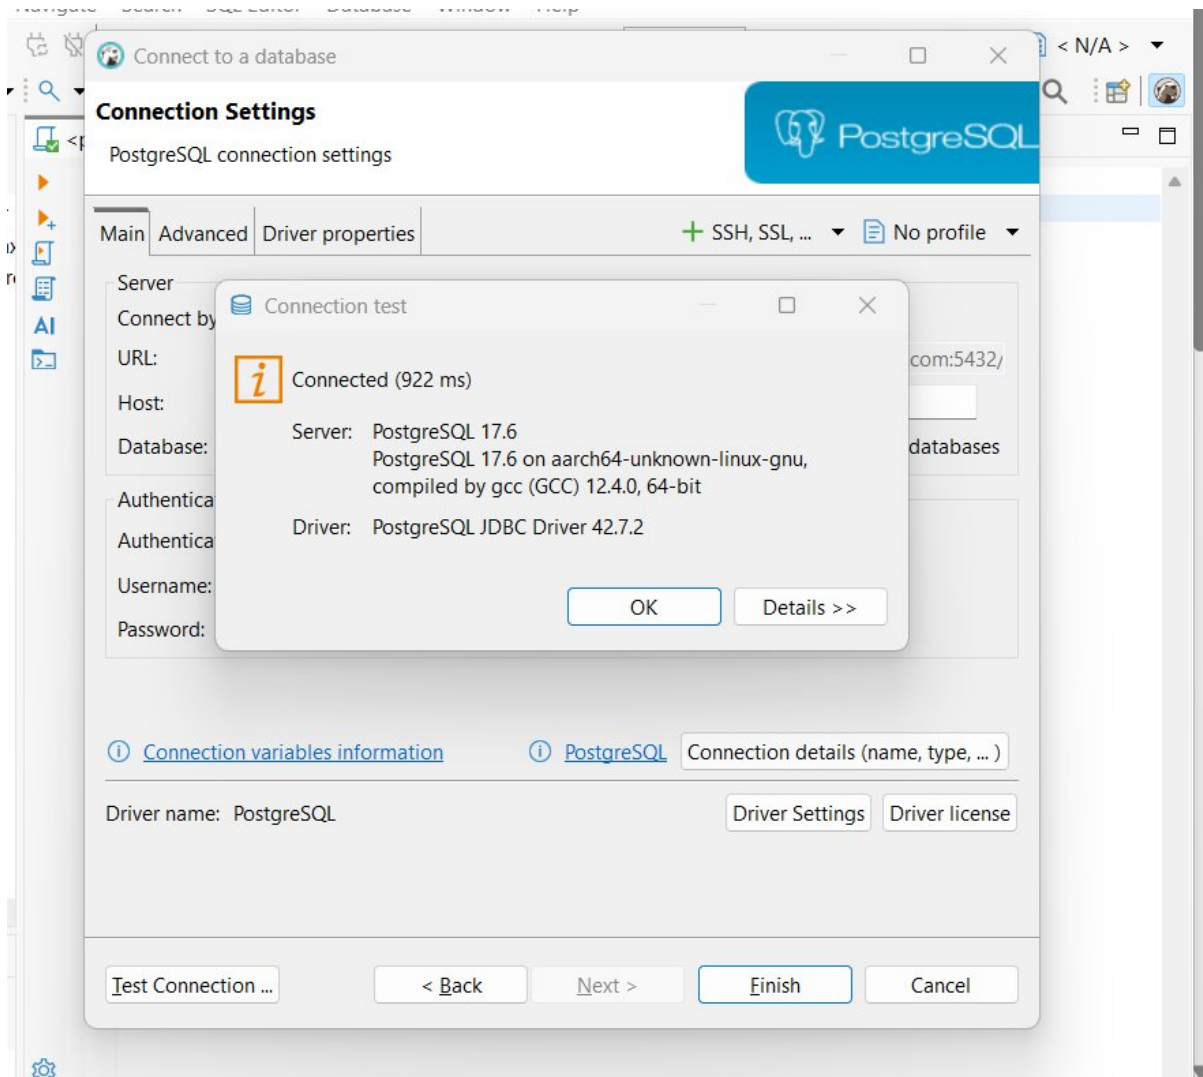
Test Connection ...

< Back

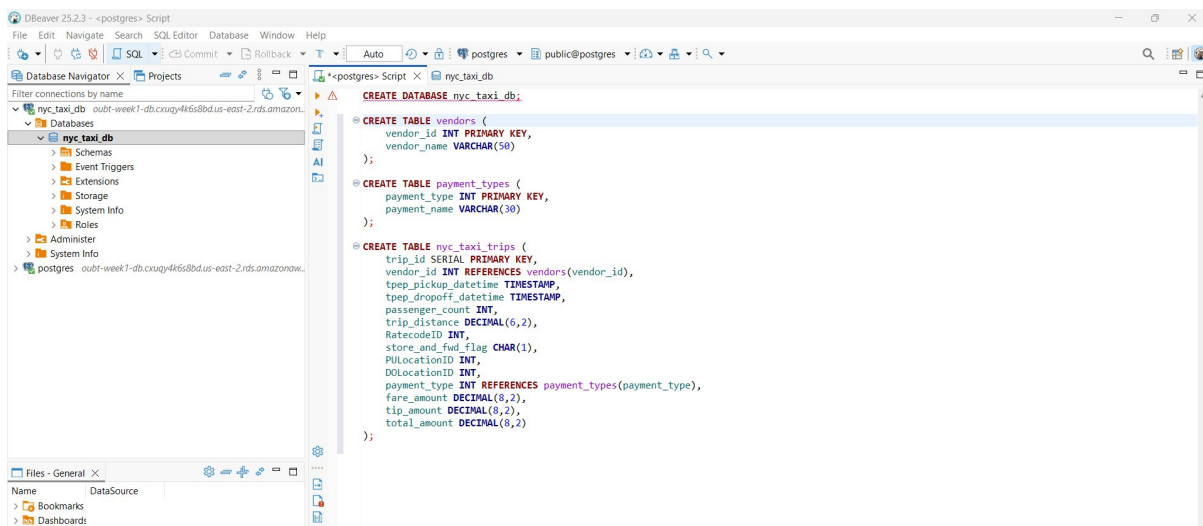
Next >

Finish

Cancel



Created a new database nyc_taxi_db.



Designed schema with three tables — vendors, payment_types, and nyc_taxi_trips — following a **Star Schema** model.

The first screenshot shows the 'vendors' table with the following data:

vendor_id	vendor_name
1	Creative Mobile Technologies
2	VeriFone Inc.
3	OUBT Transport Services
4	Likhith Mobility Labs

The second screenshot shows the 'payment_types' table with the following data:

payment_type	payment_name
1	Credit Card
2	Cash
3	Online

The third screenshot shows the 'nyc_taxi_trips' table with the following data:

trip_id	vendor_id	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance
1	2	2023-01-01 00:32:10.000	2023-01-01 00:40:36.000	1	0.97
2	2	2023-01-01 00:55:08.000	2023-01-01 01:01:27.000	1	1.1
3	2	2023-01-01 00:25:04.000	2023-01-01 00:37:49.000	1	2.51
4	1	2023-01-01 00:03:48.000	2023-01-01 00:13:25.000	0	1.9
5	2	2023-01-01 00:10:29.000	2023-01-01 00:21:19.000	1	1.43
6	2	2023-01-01 00:50:34.000	2023-01-01 01:02:52.000	1	1.84
7	2	2023-01-01 00:09:22.000	2023-01-01 00:19:49.000	1	1.66
8	2	2023-01-01 00:27:12.000	2023-01-01 00:49:56.000	1	11.7

Inserted sample trip data and reference data for vendors and payment types.

Executed multiple **SQL queries** covering JOIN, GROUP BY, aggregations, filtering, window functions, and CTEs to analyze trip records.

```
SELECT v.vendor_name,
COUNT(*) AS total_trips,
SUM(fare_amount) AS total_fare
FROM nyc_taxi_trips t
JOIN vendors v ON t.vendor_id = v.vendor_id
GROUP BY v.vendor_name;
```

```
SELECT p.payment_name,
       ROUND(AVG(fare_amount),2) AS avg_fare
FROM nyc_taxi_trips t
JOIN payment_types p ON t.payment_type = p.payment_type
GROUP BY p.payment_name;
```

```
SELECT trip_id, vendor_id, fare_amount
FROM nyc_taxi_trips
ORDER BY fare_amount DESC
LIMIT 3;
```

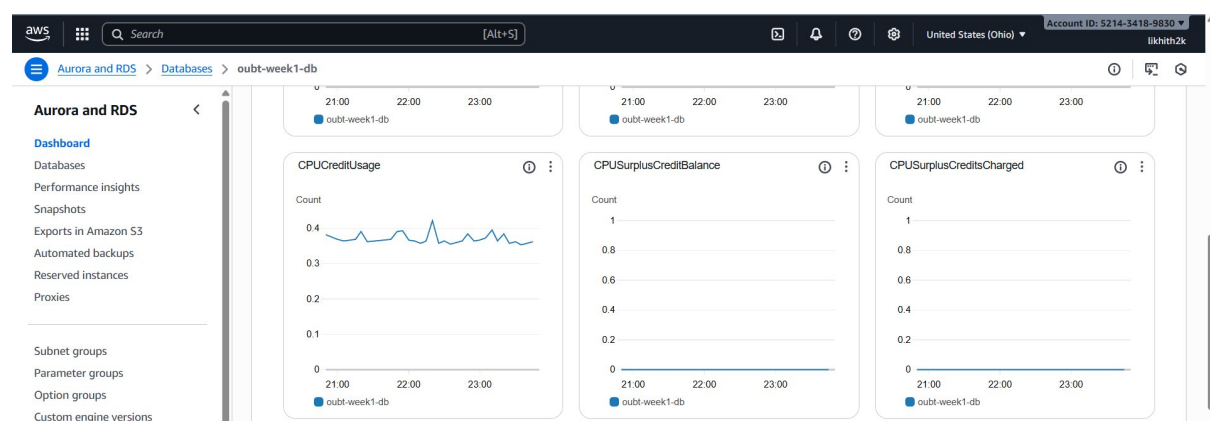
```
SELECT trip_id, vendor_id, fare_amount,
       RANK() OVER(ORDER BY fare_amount DESC) AS fare_rank
FROM nyc_taxi_trips;
```

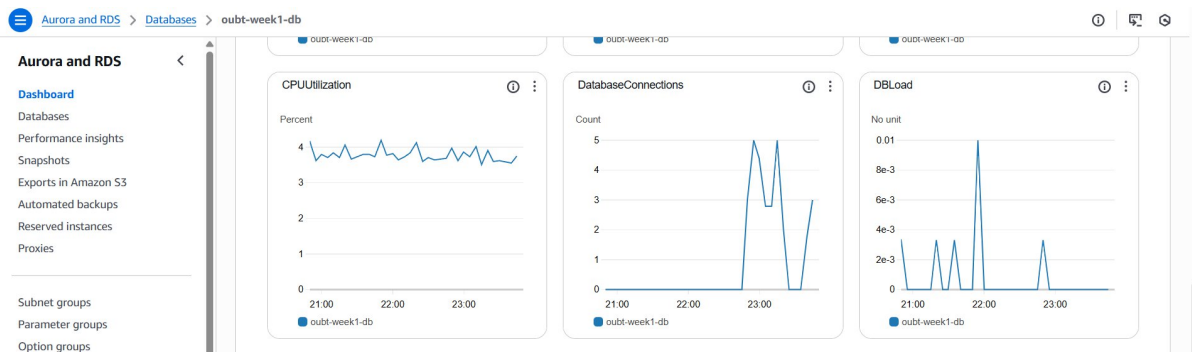
Verified query results and captured screenshots of successful execution in DBeave

CloudWatch & RDS Monitoring Test

To verify my AWS RDS instance activity, I ran multiple SQL queries from DBeaver and monitored performance in the RDS dashboard. The CloudWatch graphs showed clear spikes in **CPU utilization**, **DB load**, and **active connections** while queries were executing. This confirmed that my PostgreSQL instance was processing live workloads.

This exercise helped me understand how CloudWatch metrics can be used to monitor database health and query performance in real time.





Key Learnings:

Gained hands-on experience in **AWS RDS setup**, database connectivity, schema creation, and SQL querying through DBeaver.

Configured **IAM permissions and security groups** to enable secure database access.

Understood how **Star Schema modeling** supports analytical queries and relational design.

Practiced organizing and structuring data for better readability, consistency, and performance.